[S06] Diffusion Models: From Adding Noise to Generating Images

Vong Chi Van Le Dat Phuc An
Faculty of Information Technology, University of Science, VNU-HCM
{24125049, 24125052}@student.hcmus.edu.vn

Abstract

Diffusion models have emerged as a dominant paradigm in generative artificial intelligence, achieving unprecedented visual fidelity and robust likelihood-based training without adversarial instability. This report investigates the fundamental guiding question: *how can learning to remove noise enable a neural network to generate novel images, and why does diffusion within a learned latent space drastically reduce computational requirements?* Grounded in two foundational papers, Denoising Diffusion Probabilistic Models (Ho et al., 2020) and Latent Diffusion Models (Rombach et al., 2022), we examine the mathematical principles of forward Gaussian corruption, reverse denoising, and perceptual latent compression. Through theoretical derivations, a step-by-step numerical worked example, and empirical benchmark experiments, we provide an integrated analysis of diffusion dynamics, high-resolution scaling trade-offs, and practical deployment limitations.

Keywords: Generative AI, Denoising Diffusion Probabilistic Models, Latent Diffusion Models, Variational Autoencoders, Cross-Attention, Generative Learning Trilemma.

1 Introduction

Synthesizing high-dimensional continuous signals such as natural images is a central challenge in artificial intelligence. A natural image tensor is represented as $\mathbf{x} \in \mathbb{R}^{H \times W \times C}$, where H, W, and C denote the height, width, and color channel dimensions, respectively. Unlike supervised discriminative learning, which only seeks decision boundaries separating predefined class labels, a generative model must capture the true underlying probability distribution $p_{\text{data}}(\mathbf{x})$ of a complex, low-dimensional data manifold embedded in high-dimensional ambient space. Once learned, the model enables the synthesis of novel, realistic samples $\mathbf{x}_{\text{new}} \sim p_\theta(\mathbf{x})$.

Deep generative architectures have historically struggled to simultaneously achieve high sample sharpness, comprehensive mode diversity, and fast sampling speed. While early generative paradigms made substantial progress, each faced inherent trade-offs between sample quality, training stability, and sampling efficiency.

Tasks and Report Organization. To systematically address the guiding question in accordance with the seminar requirements, this report is structured around six concrete tasks:

- *Task 1: Foundational Paradigms and the Generative Trilemma.* We contrast discriminative classifiers with generative modeling and analyze autoregressive models, variational autoencoders, generative adversarial networks, and diffusion models across the generative learning trilemma (Sections 2.1 and 2.2).
- *Task 2: DDPM Mathematical Framework.* We formulate the variance-preserving forward process, variance schedule selection, closed-form multi-step transitions, the learned reverse Markov chain, and the simplified noise-prediction training objective (Section 3.1).
- *Task 3: Symbol Definitions and 1D Worked Example.* We specify complete symbol definitions, physical units, and modeling assumptions, accompanied by a step-by-step one-dimensional numerical worked example illustrating stochastic signal contraction and reverse denoising (Section 3.2).
- *Task 4: Multi-Level Forward Corruption and Reverse Image Generation.* We demonstrate forward noise corruption across multiple timesteps and iterative reverse image generation trajectories on empirical benchmarks (Section 4.2).
- *Task 5: Latent Diffusion Architecture.* We analyze and empirically validate the two-stage inductive bias of Latent Diffusion Models, examining perceptual autoencoder compression, latent manifold smoothness, and cross-attention conditioning with classifier-free guidance (Sections 3.3 and 4.3).

- *Task 6: Pixel vs. Latent Comparative Evaluation and Practical Limitations.* We provide an evidence-supported comparison between pixel-space and latent-space diffusion across explicit computational and quality metrics, followed by a critical evaluation of real-world limitations including sequential sampling latency, prompt attribute binding, copyright memorization, and societal bias (Sections 4.4 and 5).

2 Background and Paper Review

2.1 Distinguishing Discriminative vs. Generative Models

Supervised machine learning and probabilistic modeling are fundamentally divided into two complementary paradigms:

1. *Discriminative Models:* Discriminator models are mainly used for classification and regression tasks. A discriminative model directly learns to predict the conditional probability distribution $P(Y \mid \mathbf{x})$, where Y is the class label and $\mathbf{x}$ represents the input features. It learns decision boundaries separating classes without modeling the distribution of the input features $\mathbf{x}$. Consequently, a discriminative model cannot generate novel data samples and often produces spuriously overconfident predictions on empty out-of-distribution regions.

2. *Generative Models:* Generative models are primarily used to generate novel data, though they can also be used for classification. A generative model learns the underlying data distribution $p(\mathbf{x})$, or the joint distribution $p(\mathbf{x}, Y)$ when conditioned on class labels Y. Its primary function is synthesizing new, realistic samples $\mathbf{x}_{\text{new}} \sim p(\mathbf{x})$ that resemble the training data distribution, while assigning negligible density to regions devoid of data.

Figure 1 illustrates this fundamental distinction on a synthetic four-class two-dimensional manifold. In Figure 1a, the discriminative classifier learns quadrant decision boundaries to predict $P(Y \mid \mathbf{x})$; however, when queried at an empty out-of-distribution point $\mathbf{x}_{\text{OOD}} = (4.5, 4.5)$, it extrapolates with unwarranted 100.0% overconfidence and lacks any mechanism to synthesize new observations. In contrast, Figure 1b demonstrates the generative paradigm: the model explicitly learns the data density $p_\theta(\mathbf{x})$, reliably identifies out-of-distribution queries by assigning near-zero density ($p(\mathbf{x}_{\text{OOD}}) \approx 0$), and generates novel realistic samples $\mathbf{x}_{\text{new}} \sim p_\theta(\mathbf{x})$ across all four data modes.

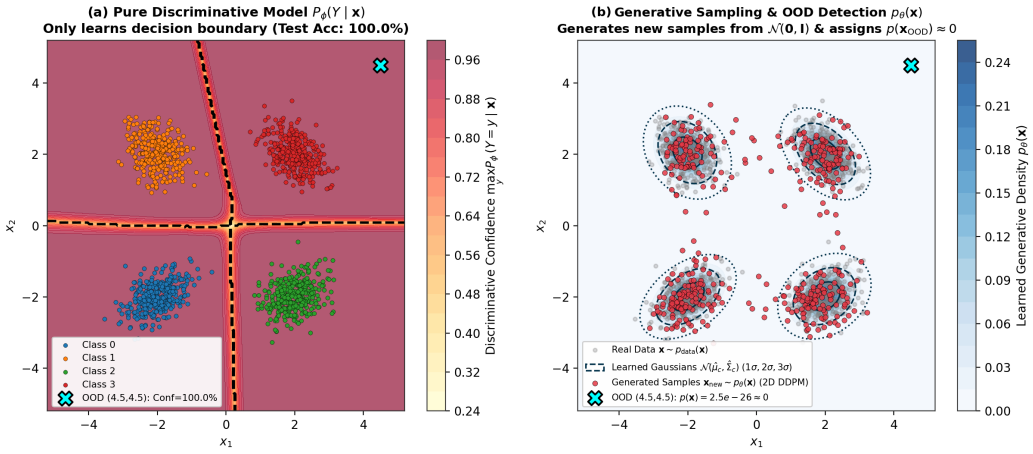


Figure 1. Visual comparison of modeling paradigms on a four-class two-dimensional manifold (Task 1): (a) A discriminative classifier learns quadrant decision boundaries to predict $P(Y \mid \mathbf{x})$ for classification, but exhibits unwarranted 100.0% overconfidence at an empty out-of-distribution coordinate $\mathbf{x}_{\text{OOD}} = (4.5, 4.5)$ and lacks generative ability; (b) A generative model explicitly learns the data density $p_\theta(\mathbf{x})$, detects out-of-distribution queries by assigning negligible density ($p(\mathbf{x}_{\text{OOD}}) \approx 0$), and synthesizes novel realistic samples $\mathbf{x}_{\text{new}} \sim p_\theta(\mathbf{x})$ along all four data modes.

2.2 Qualitative comparison of GANs, autoregressive models, and diffusion models

To contextualize diffusion models within the generative literature, we compare autoregressive models, variational autoencoders (VAEs), generative adversarial networks (GANs), and diffusion models across the Generative Learning Trilemma [16], which posits that a generative model ideally satisfies three core properties: (1) *High Sample Quality* (generating sharp, photorealistic samples), (2) *Fast Sampling* (synthesizing samples with minimal forward passes), and (3) *Full Mode Coverage* (faithfully covering all modes of the data distribution without mode collapse).

Historically, generative architectures achieved at most two of these three objectives, a theoretical trade-off that we also verify directly in our empirical experiments (Sections 4.2–4.4):

1. *Autoregressive Models (PixelRNN, PixelCNN) [10, 11]:* Factorize the joint distribution pixel-by-pixel via the probability chain rule $p_\theta(\mathbf{x}) = \prod_{i=1}^{n} p_\theta(x^{(i)} \mid x^{(<i)})$. They maximize exact log-likelihood and avoid mode collapse, achieving high quality and full diversity. However, sequential pixel-by-pixel generation incurs $\mathcal{O}(n)$ complexity and extremely slow sampling, as confirmed in our MNIST benchmark (Section 4.2, Figure 7), where PixelCNN requires 784 sequential forward passes ($4{,}294.7$ ms per 64 images on an NVIDIA Tesla T4 GPU).

2. *Variational Autoencoders (VAEs) [4]:* Introduce a continuous latent variable $\mathbf{z} \sim p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I})$ with an encoder $q_\phi(\mathbf{z} \mid \mathbf{x})$ and decoder $p_\theta(\mathbf{x} \mid \mathbf{z})$ trained via the evidence lower bound (ELBO). VAEs generate images in a single decoder pass and cover all modes, but pixel-level reconstruction losses average across plausible modes and produce visibly blurry samples, reflected in our benchmark by the fastest T4 GPU sampling latency (6.5 ms per 64 images) but the lowest edge sharpness (LapVar $= 0.0182$, FID $= 82.27$ in Table 3).

3. *Generative Adversarial Networks (GANs) [5, 6]:* Frame synthesis as a minimax game $\min_G \max_D$ between a Generator $G_\theta(\mathbf{z})$ and Discriminator $D_\phi(\mathbf{x})$. GANs produce razor-sharp images in a single feed-forward step (11.4 ms per 64 images and LapVar $= 0.0541$ in Table 3), but the lack of an explicit density objective causes training instability and mode dropping (FID $= 44.10$) [8].

4. *Denoising Diffusion Probabilistic Models (DDPM) [3, 1]:* Invert a parameterized T-step Markov chain using a U-Net $\boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)$ trained via unweighted mean-squared error L_{simple}. Diffusion models achieve state-of-the-art visual fidelity and full mode coverage, attaining the highest sharpness (LapVar $= 0.0601$), highest classifier confidence (97.4%), and lowest FID (16.85) in our MNIST benchmark (Table 3), with sequential sampling latency ($3{,}926.8$ ms for $T = 300$ steps on Tesla T4 GPU) as their primary bottleneck.

5. *Latent Diffusion Models (LDM) [2]:* Resolve the computational bottleneck of DDPM by compressing $\mathbf{x}_0 \to \mathbf{z}_0$ via a frozen perceptual spatial VAE and executing diffusion on the low-dimensional latent tensor $\mathbf{z}_t$. Our experiments in Sections 4.3 and 4.4 (Table 4) demonstrate that this spatial latent bottleneck preserves perceptual fidelity (PSNR $= 33.74$ dB, FID $= 0.18$) while accelerating U-Net training by $7.61\times$ per epoch and cutting GPU activation memory by $13.81\times$ on a Tesla T4 GPU.

Table 1 summarizes the qualitative properties of each generative family across the trilemma, citing primary literature sources for every claim, while Section 4.2 (Figure 7 and Table 3) and Sections 4.3–4.4 (Table 4) present our corresponding qualitative and quantitative experimental validations.

Table 1. Qualitative comparison of major generative model families across the Generative Learning Trilemma [16], citing primary literature sources for each theoretical property and empirical claim (Task 1).

Model Family	Density $p_\theta(\mathbf{x})$	Sample Quality	Mode Coverage	Sampling Speed
Autoregressive [10, 11] (PixelRNN, PixelCNN)	Explicit, Exact (Exact MLE) [10, 11]	Sharp local contrast [10, 16]	Full (No mode collapse; stable MLE) [10, 1]	Extremely Slow ($\mathcal{O}(n)$ sequential pixel evaluations) [10, 13, 16]
VAE [4] (Variational Autoencoder)	Explicit, Approx. (ELBO Bound) [4]	Blurry (Single-step L_2/BCE mode averaging) [4, 9, 16]	Full, smooth latent manifold [4, 16]	Very Fast (1 Decoder forward pass) [4, 16]
GAN [5, 6] (DCGAN, StyleGAN)	Implicit (Likelihood-free) [5]	Razor Sharp (Mode-seeking via adversarial loss) [5, 7, 16]	Poor (Prone to Mode Collapse and minimax instability) [5, 8, 16]	Very Fast (1 Generator forward pass) [5, 16]
Pixel Diffusion [3, 1] (DDPM)	Explicit, Approx. (T-step ELBO) [3, 1]	SOTA Sharpness (FID 3.17 on CIFAR-10) [1, 14]	Full (Stable regression; beats GAN mode coverage) [1, 14, 16]	Slow ($T \approx 1000$ sequential full-res U-Net passes) [1, 13, 16]
Latent Diffusion [2] (LDM / Stable Diffusion)	Two-Stage ELBO (VAE + Diffusion) [2]	SOTA Photorealism at high resolution (FID 5.11 on CelebA-HQ) [2]	Full and flexible multimodal text conditioning [2]	Fast and Practical (Operates on $48\times$ smaller $\mathbf{z}_t$; 20 to 50 DDIM steps) [2, 13]

2.3 Foundational Literature Review

This report analyzes two foundational milestone papers:

- *Required Paper: Ho, Jain and Abbeel (NeurIPS 2020), "Denoising Diffusion Probabilistic Models" (DDPM) [1]:* Demonstrated that diffusion models generate image quality rivaling GANs while retaining likelihood-based mode coverage. The paper made three key contributions: (1) establishing the equivalence between denoising score matching and Markov chain reverse step estimation; (2) reparameterizing the reverse transition mean to predict the injected Gaussian noise vector $\boldsymbol{\epsilon}$; (3) introducing the unweighted simplified training objective $L_{\text{simple}}(\theta)$, which suppresses small-step noise artifacts and prioritizes perceptual visual features.

- *Recommended Paper: Rombach et al. (CVPR 2022), "High-Resolution Image Synthesis With Latent Diffusion Models" (LDM) [2]:* Identified that training pixel-space diffusion models on high-resolution images allocates excessive compute to imperceptible high-frequency details. Rombach et al. introduced a two-stage paradigm: (1) a first stage Variational Autoencoder (VAE) trained with perceptual and adversarial losses that compresses high-dimensional images into an efficient spatial latent space; (2) a second stage diffusion backbone operating exclusively on latent representations, incorporating cross-attention layers to enable flexible conditioning on text prompts, bounding boxes, or segmentation maps.

3 Method and Step-by-Step Worked Example

3.1 Forward Process, Variance Schedule, Reverse Denoising, and Sampling (Task 2)

Forward Noising Process. Let clean data follow an unknown distribution $\mathbf{x}_0 \sim q(\mathbf{x}_0)$. The forward diffusion process is a non-trainable Markov chain governed by a fixed variance schedule $\beta_1, \beta_2, \ldots, \beta_T \in (0, 1)$, where the joint distribution and individual Gaussian transitions are defined as:

$$q(\mathbf{x}_{1:T} \mid \mathbf{x}_0) = \prod_{t=1}^{T} q(\mathbf{x}_t \mid \mathbf{x}_{t-1}), \quad q(\mathbf{x}_t \mid \mathbf{x}_{t-1}) = \mathcal{N}\big(\mathbf{x}_t;\ \sqrt{1-\beta_t}\,\mathbf{x}_{t-1},\ \beta_t\mathbf{I}\big) \tag{1}$$

Variance-Preserving Property. Equation (1) implements a Variance-Preserving (VP) stochastic process. By scaling the incoming signal by $\sqrt{1-\beta_t}$, total variance is preserved: if the previous state has unit variance $\mathrm{Var}(\mathbf{x}_{t-1}) = \mathbf{I}$, then $\mathrm{Var}(\mathbf{x}_t) = (\sqrt{1-\beta_t})^2\mathrm{Var}(\mathbf{x}_{t-1}) + \beta_t\mathbf{I} = (1-\beta_t)\mathbf{I} + \beta_t\mathbf{I} = \mathbf{I}$. Consequently, the signal amplitude never explodes, and the final state converges smoothly to standard white Gaussian noise $\mathbf{x}_T \approx \mathcal{N}(\mathbf{0}, \mathbf{I})$.

Closed-Form Multi-Step Transition. A key property of Gaussian Markov chains is that intermediate steps can be bypassed in closed form. Defining $\alpha_t = 1 - \beta_t$ and cumulative signal retention $\bar{\alpha}_t = \prod_{s=1}^{t} \alpha_s$, any intermediate noisy state $\mathbf{x}_t$ can be sampled directly from $\mathbf{x}_0$ using standard Gaussian noise $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$:

$$q(\mathbf{x}_t \mid \mathbf{x}_0) = \mathcal{N}\big(\mathbf{x}_t;\ \sqrt{\bar{\alpha}_t}\,\mathbf{x}_0,\ (1-\bar{\alpha}_t)\mathbf{I}\big) \iff \mathbf{x}_t = \sqrt{\bar{\alpha}_t}\,\mathbf{x}_0 + \sqrt{1-\bar{\alpha}_t}\,\boldsymbol{\epsilon} \tag{2}$$

Equation (2) enables efficient training by sampling random timesteps $t \sim \mathrm{Uniform}(\{1, \ldots, T\})$ in $\mathcal{O}(1)$ time without simulating steps 1 to $t-1$.

Variance Schedule Selection. The schedule $\beta_1, \ldots, \beta_T$ controls the rate of information destruction:

- *Small β_t and Large T:* Small β_t ensures that the true reverse transition $q(\mathbf{x}_{t-1} \mid \mathbf{x}_t)$ remains nearly Gaussian, while sufficiently large T ensures $\bar{\alpha}_T \approx 0$ so that $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$.
- *Linear Schedule (Ho et al. [1]):* Linearly interpolates from $\beta_1 = 10^{-4}$ to $\beta_T = 0.02$ across $T = 1000$.
- *Cosine Schedule (Nichol & Dhariwal [15]):* Sets $\bar{\alpha}_t = f(t)/f(0)$ with $f(t) = \cos^2\big(\frac{t/T+s}{1+s} \cdot \frac{\pi}{2}\big)$ to prevent abrupt signal destruction in early timesteps.
- *Latent Diffusion Schedule (Rombach et al. [2]):* Uses a scaled-linear schedule ($\beta_1 = 0.00085 \to \beta_T = 0.012$) tailored to normalized latent spaces.

Reverse Denoising Process and Posterior Tractability. The marginal reverse transition $q(\mathbf{x}_{t-1} \mid \mathbf{x}_t)$ is intractable because it requires integrating over the unknown data distribution $q(\mathbf{x}_0)$. However, conditioning on the clean image $\mathbf{x}_0$ via Bayes' rule renders the forward posterior Gaussian and analytically tractable:

$$q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0) = \frac{q(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{x}_0)\, q(\mathbf{x}_{t-1} \mid \mathbf{x}_0)}{q(\mathbf{x}_t \mid \mathbf{x}_0)} = \mathcal{N}\big(\mathbf{x}_{t-1};\ \tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0),\ \tilde{\beta}_t\mathbf{I}\big) \tag{3}$$

Matching quadratic Gaussian exponents yields the closed-form posterior mean and variance:

$$\tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0) = \frac{\sqrt{\bar{\alpha}_{t-1}}\,\beta_t}{1-\bar{\alpha}_t}\,\mathbf{x}_0 + \frac{\sqrt{\alpha_t}\,(1-\bar{\alpha}_{t-1})}{1-\bar{\alpha}_t}\,\mathbf{x}_t, \quad \tilde{\beta}_t = \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\,\beta_t \tag{4}$$

Since $\mathbf{x}_0$ is unknown during inference, we train a parameterized transition $p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t) = \mathcal{N}\big(\mathbf{x}_{t-1};\ \boldsymbol{\mu}_\theta(\mathbf{x}_t, t),\ \sigma_t^2\mathbf{I}\big)$ with $\sigma_t^2 = \tilde{\beta}_t$ (or β_t). Rearranging Equation (2) as $\mathbf{x}_0 = \frac{1}{\sqrt{\bar{\alpha}_t}}\big(\mathbf{x}_t - \sqrt{1-\bar{\alpha}_t}\,\boldsymbol{\epsilon}\big)$ and substituting into Equation (4) reparameterizes $\boldsymbol{\mu}_\theta$ in terms of a U-Net noise predictor $\boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)$, yielding the reverse sampling update step from $\mathbf{x}_t$ to $\mathbf{x}_{t-1}$:

$$\boldsymbol{\mu}_\theta(\mathbf{x}_t, t) = \frac{1}{\sqrt{\alpha_t}}\left(\mathbf{x}_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}}\,\boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)\right) \implies \mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}}\left(\mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}}\,\boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)\right) + \sigma_t\mathbf{z} \tag{5}$$

where $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ for $t > 1$, and $\mathbf{z} = \mathbf{0}$ at the final step $t = 1$.

Score-Matching Equivalence and Simplified Objective L_{simple}. By Tweedie's Formula and denoising score matching [12, 1], the predicted noise estimates the Stein score of the perturbed density, $\mathbf{s}_\theta(\mathbf{x}_t, t) = -\boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)/\sqrt{1-\bar{\alpha}_t} \approx \nabla_{\mathbf{x}_t} \log q_t(\mathbf{x}_t)$, so subtracting predicted noise performs a step toward higher data density. Optimizing the variational lower bound L_{VLB} over the KL divergence terms $L_{t-1} = D_{\text{KL}}\big(q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0) \parallel p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t)\big)$ reduces to a weighted MSE on $\boldsymbol{\epsilon}$. Ho et al. [1] showed that dropping the timestep weighting prefactor improves sample quality by emphasizing perceptually salient larger-t denoising tasks, yielding the simplified objective:

$$L_{\text{simple}}(\theta) = \mathbb{E}_{t,\mathbf{x}_0,\boldsymbol{\epsilon}}\left[\left\|\boldsymbol{\epsilon} - \boldsymbol{\epsilon}_\theta\big(\sqrt{\bar{\alpha}_t}\,\mathbf{x}_0 + \sqrt{1-\bar{\alpha}_t}\,\boldsymbol{\epsilon},\, t\big)\right\|^2\right] \tag{6}$$

where $t \sim \text{Uniform}(\{1, \ldots, T\})$, $\mathbf{x}_0 \sim q(\mathbf{x}_0)$, and $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$. Figure 2 and Figure 3 illustrate the Markov chain and the complete DDPM training and sampling algorithms.

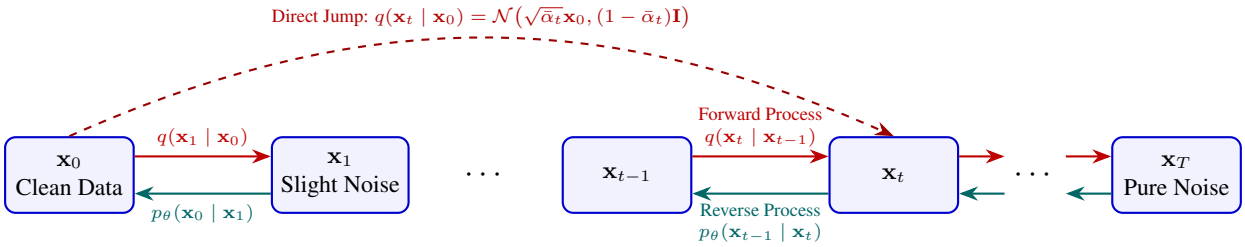


Figure 2. The DDPM Markov chain architecture: forward noising chain $q(\mathbf{x}_t \mid \mathbf{x}_{t-1})$ (top red arrows), closed-form multi-step jump $q(\mathbf{x}_t \mid \mathbf{x}_0)$ (curved red dashed arrow), and learned reverse denoising chain $p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t)$ parameterized by U-Net $\boldsymbol{\epsilon}_\theta$ (bottom teal arrows) [1].

Algorithm 1 Training

repeat
  $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
  $t \sim \text{Uniform}(\{1, \ldots, T\})$
  $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
  Take gradient descent step on
  $\nabla_\theta \left\|\boldsymbol{\epsilon} - \boldsymbol{\epsilon}_\theta\big(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1-\bar{\alpha}_t}\boldsymbol{\epsilon},\, t\big)\right\|^2$
until converged

Algorithm 2 Sampling

$\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
for $t = T, \ldots, 1$ **do**
  $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ **if** $t > 1$ **else** $\mathbf{z} = \mathbf{0}$
  $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}}\left(\mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}}\boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)\right) + \sigma_t \mathbf{z}$
end for
return $\mathbf{x}_0$

Figure 3. DDPM Training (Algorithm 1) and Sampling (Algorithm 2) matching Ho et al. [1].

3.2 Formal Equation Specifications and 1D Worked Numerical Example (Section III & Task 3)

Symbol Definitions and Modeling Assumptions. Table 2 provides complete mathematical definitions, units, value ranges, and explicit modeling assumptions for every symbol in the DDPM framework.

Step-by-Step One-Dimensional Worked Example. To verify the equations concretely, consider a one-dimensional continuous coordinate ($D = 1$) with initial clean signal $x_0 = 0.80$ at timestep $t = 1$, forward variance parameter $\beta_1 = 0.19$ (so $\alpha_1 = \bar{\alpha}_1 = 1 - \beta_1 = 0.81$ and $\sqrt{\alpha_1} = 0.90$), and sampled standard Gaussian noise $\epsilon = -0.50 \sim \mathcal{N}(0, 1)$:

- *Step 1: Forward Distribution $q(x_1 \mid x_0 = 0.80)$:* From Equation (1), $q(x_1 \mid x_0) = \mathcal{N}\big(x_1; \sqrt{1-\beta_1}\,x_0, \beta_1\big)$. Substituting the parameters yields mean $\mu_{x_1} = \sqrt{1-\beta_1}\,x_0 = 0.90 \times 0.80 = 0.7200$ and standard deviation $\sigma_{x_1} = \sqrt{\beta_1} = \sqrt{0.19} \approx 0.4359$, so $q(x_1 \mid x_0 = 0.80) = \mathcal{N}(x_1;\, 0.7200,\, 0.19)$.
- *Step 2: Realized Noisy Pixel x_1 for $\epsilon = -0.50$:* Applying the reparameterization formula in Equation (2):

$$x_1 = \sqrt{\bar{\alpha}_1}\,x_0 + \sqrt{1-\bar{\alpha}_1}\,\epsilon = 0.90(0.80) + 0.4359(-0.50) = 0.5021 \tag{7}$$

- *Step 3: Physical Interpretation of Stochastic Expectation $\mathbb{E}[x_1 \mid x_0] = 0.7200$:* Because $\mathbb{E}[\epsilon] = 0$, the expected corrupted coordinate is $\mathbb{E}[x_1 \mid x_0] = \sqrt{1-\beta_1}\,x_0 + \sqrt{\beta_1}\,\mathbb{E}[\epsilon] = 0.90(0.80) + 0 = 0.7200$. The factor $\sqrt{1-\beta_1} = 0.90 < 1.0$ acts as a deterministic signal contraction toward the origin 0 that balances the injected variance $\beta_1 = 0.19$ to preserve unit total variance.

Table 2. Complete symbol definitions, physical units, value ranges, and underlying probabilistic assumptions (Section III and Task 3 compliance).

Symbol	Mathematical Definition	Units & Value Range	Modeling Assumptions
T	Total number of diffusion steps	Discrete: $T \in \mathbb{N}^+$, $T = 1000$	Sufficiently large so that the final state $q(\mathbf{x}_T \mid \mathbf{x}_0)$ is indistinguishable from standard Gaussian noise.
t	Discrete diffusion timestep	Discrete: $t \in \{1, 2, \ldots, T\}$	Sampled uniformly during training ($t \sim \text{Uniform}(\{1, \ldots, T\})$) to ensure balanced multiscale noise feature learning.
$\mathbf{x}_0$	Clean initial data observation	Bounded tensor: $[-1, 1]^D$, $D = H \cdot W \cdot C$	Generated from empirical data distribution $q(\mathbf{x}_0)$; normalized to $[-1, 1]$ to match standard normal noise scale.
$\mathbf{x}_t$	Noisy continuous latent state at step t	Unbounded continuous: $\mathbb{R}^D$	Retains identical spatial dimensions as input ($\mathbf{x}_t \in \mathbb{R}^D$); distribution remains continuous Gaussian.
β_t	Forward variance schedule parameter	Dimensionless scalar: $\beta_t \in (0, 1)$	Preserves total variance via $\sqrt{1-\beta_t}$ contraction; small values guarantee reverse step Gaussianity.
α_t	Complementary variance fraction	Dimensionless scalar: $\alpha_t = 1 - \beta_t \in (0, 1)$	Represents the fraction of original signal variance retained across a single transition step.
$\bar{\alpha}_t$	Cumulative signal retention factor	Dimensionless scalar: $\bar{\alpha}_t = \prod_{s=1}^{t} \alpha_s \in (0, 1)$	Strictly monotonically decreasing toward zero ($\bar{\alpha}_1 \approx 0.9999 \to \bar{\alpha}_T \approx 0$).
$\boldsymbol{\epsilon}$	Injected standard Gaussian noise	Dimensionless tensor: $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_D)$	Independent and identically distributed across all dimensions and timesteps: $\mathbb{E}[\boldsymbol{\epsilon}] = \mathbf{0}$, $\text{Cov}(\boldsymbol{\epsilon}) = \mathbf{I}$.
$\boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)$	Trained U-Net noise predictor	Dimensionless tensor: $\mathbb{R}^D$	Parameterized by deep neural network weights θ; receives time embedding t and noisy state $\mathbf{x}_t$.
σ_t^2	Reverse stochastic variance	Dimensionless scalar: $\sigma_t^2 \in [0, \beta_t]$	Set to posterior variance $\tilde{\beta}_t = \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\beta_t$ or upper bound β_t [1].
$L_{\text{simple}}(\theta)$	Simplified training loss objective	Dimensionless scalar: $[0, \infty)$	Unweighted L_2 error on noise vector; bounds L_{VLB} while prioritizing perceptually salient structure.

- *Step 4: Worked Calculation of L_{simple} and Reverse Denoising ($x_1 \to \hat{x}_0$):* Suppose the neural network evaluates ($x_1 = 0.5021$, $t = 1$) and predicts $\epsilon_\theta(0.5021, 1) = -0.30$. The single-sample squared error loss is $L_{\text{simple}} = (\epsilon - \epsilon_\theta(x_1, 1))^2 = (-0.50 - (-0.30))^2 = 0.0400$. At inference ($t = 1$, where $\sigma_1 = 0$), Equation (5) denoises x_1 as:

$$\hat{x}_0 = \frac{1}{\sqrt{\alpha_1}}\left(x_1 - \frac{1-\alpha_1}{\sqrt{1-\bar{\alpha}_1}}\,\epsilon_\theta(x_1, 1)\right) = \frac{1}{0.90}\left(0.5021 - \frac{0.19}{0.4359}(-0.30)\right) = 0.7031 \tag{8}$$

  If the network were perfectly trained ($\epsilon_\theta = \epsilon = -0.50$), the reverse step yields $\hat{x}_0 = \frac{1}{0.90}\big(0.5021 - 0.4359(-0.50)\big) = 0.8000$, recovering the exact clean coordinate.

3.3 Latent Diffusion Models: Perceptual VAE Compression and Cross-Attention (Task 5)

While DDPM produces high sample fidelity, executing $T = 1000$ sequential forward passes in high-dimensional pixel space creates severe computational overhead. For example, a $512 \times 512 \times 3$ image has 786,432 dimensions. Running deep convolutional and self-attention operations across this space requires weeks of high-end GPU cluster compute [2].

Two-Stage Separation of Inductive Biases. Rombach et al. [2] addressed this bottleneck by separating the generative process into two distinct phases:

1. *Perceptual Compression Stage:* A Variational Autoencoder (VAE) trained with perceptual loss and patch-based adversarial loss compresses the RGB image $\mathbf{x} \in \mathbb{R}^{H \times W \times 3}$ into a spatial latent representation $\mathbf{z} = \mathcal{E}(\mathbf{x}) \in \mathbb{R}^{h \times w \times c}$, where $h = H/f$, $w = W/f$, and $f = 2^m$ is the spatial downsampling factor (typically $f \in \{4, 8\}$). For a $512 \times 512 \times 3$ image with $f = 8$ and $c = 4$, the latent tensor has shape $64 \times 64 \times 4 = 16{,}384$ entries ($48\times$ smaller volume). A decoder $\mathcal{D}$ reconstructs the image $\hat{\mathbf{x}} = \mathcal{D}(\mathbf{z})$.
2. *Latent Diffusion Stage:* The autoencoder parameters $\mathcal{E}$ and $\mathcal{D}$ are frozen, and a time-conditioned U-Net operates exclusively within the latent space $\mathbf{z}_0 \leftrightarrow \mathbf{z}_T$, focusing capacity on semantic structure rather than imperceptible pixel noise.

Cross-Attention Conditioning and Classifier-Free Guidance (CFG). To condition generation on flexible multimodal inputs y (such as text prompts or class labels), LDMs augment the U-Net backbone with cross-attention layers. A domain-specific encoder $\tau_\theta(y) \in \mathbb{R}^{M \times d_\tau}$ projects the prompt y, and intermediate U-Net features $\varphi_i(\mathbf{z}_t) \in \mathbb{R}^{N \times d_\epsilon}$ form queries, keys, and values $\mathbf{Q} = \mathbf{W}_Q^{(i)}\,\varphi_i(\mathbf{z}_t)$, $\mathbf{K} = \mathbf{W}_K^{(i)}\,\tau_\theta(y)$, and $\mathbf{V} = \mathbf{W}_V^{(i)}\,\tau_\theta(y)$:

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d}}\right)\mathbf{V}, \quad L_{\text{LDM}}(\theta) = \mathbb{E}_{\mathcal{E}(\mathbf{x}_0),\, y,\, \boldsymbol{\epsilon},\, t}\left[\left\|\boldsymbol{\epsilon} - \boldsymbol{\epsilon}_\theta\big(\mathbf{z}_t,\, t,\, \tau_\theta(y)\big)\right\|^2\right] \tag{9}$$

where $\mathbf{z}_t = \sqrt{\bar{\alpha}_t}\,\mathcal{E}(\mathbf{x}_0) + \sqrt{1-\bar{\alpha}_t}\,\boldsymbol{\epsilon}$. During training, the conditioning token is randomly replaced by a null token $\varnothing$ (e.g., $p_{\text{uncond}} = 0.15$) so that at inference Classifier-Free Guidance (CFG) extrapolates toward the conditional direction with scale $\alpha \geq 1$:

$$\tilde{\boldsymbol{\epsilon}}_\theta(\mathbf{z}_t, t, y, \alpha) = \boldsymbol{\epsilon}_\theta(\mathbf{z}_t, t, \varnothing) + \alpha\big(\boldsymbol{\epsilon}_\theta(\mathbf{z}_t, t, \tau_\theta(y)) - \boldsymbol{\epsilon}_\theta(\mathbf{z}_t, t, \varnothing)\big) \tag{10}$$

Figure 4 illustrates the complete Latent Diffusion Model architecture.

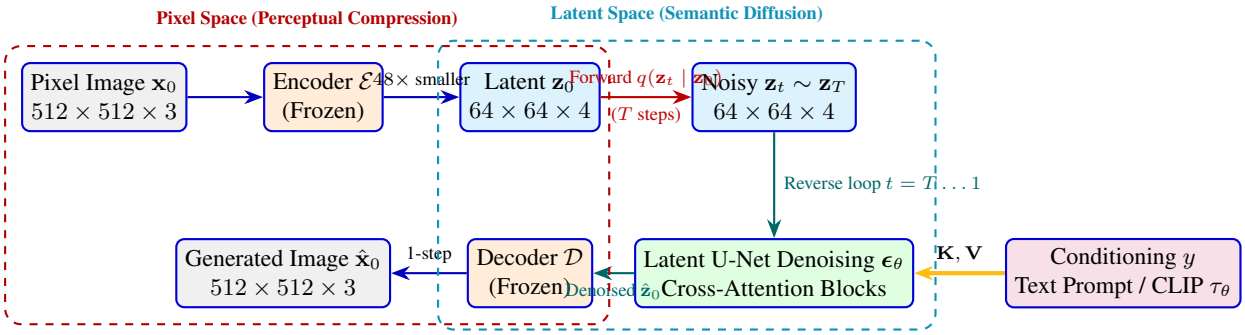


Figure 4. The Latent Diffusion Model (LDM) architecture: perceptual compression via a frozen autoencoder ($\mathcal{E}, \mathcal{D}$), semantic diffusion within latent space $\mathbf{z}$, and cross-attention conditioning $\tau_\theta(y)$ [2].

4 Results and Discussion

4.1 Implementation Setup and Datasets

To empirically validate the theoretical principles of diffusion models, our team developed and executed self-contained PyTorch experimental suites on an NVIDIA Tesla T4 GPU (14.6 GB VRAM, CUDA 13.0) across three experimental settings:

1. *Pixel-Space MNIST Generative Benchmark (Tasks 1, 4):* Trained Conv-VAE (12 epochs, 51.1 s), Autoregressive PixelCNN (12 epochs, 168.9 s), DCGAN (14 epochs, 95.7 s), and a Pixel-Space DDPM ($T = 300$, linear schedule $\beta_1 = 10^{-4} \to \beta_{300} = 0.055$, Sinusoidal time-embedded U-Net, 15 epochs, 272.3 s) on MNIST (28×28), evaluated by a trained MNIST CNN classifier (98.85% test accuracy, trained in 5.6 s on Tesla T4 GPU).

2. *Four Core Latent Diffusion Experiments (Tasks 5, 6):* Trained a Stage-1 Spatial Perceptual KL-VAE ($32 \times 32 \times 1 \to 8 \times 8 \times 4$, downsampling factor $f = 4$, $\beta_{\text{KL}} = 10^{-4}$), a 1D-Vector Bottleneck VAE ($\mathbf{z} \in \mathbb{R}^{10}$), two architecturally identical U-Nets on 32×32 pixels vs. 8×8 latents, and a Multi-Head Cross-Attention Latent U-Net conditioned on digit class prompts with Classifier-Free Guidance ($p_{\text{uncond}} = 0.15$), evaluated by a 32×32 CNN classifier (98.10% test accuracy) and pretrained Stable Diffusion KL-VAE (`stabilityai/sd-vae-ft-mse`) on CIFAR-10 RGB.

3. *Pretrained High-Resolution Checkpoints (Tasks 1, 4, 6):* Evaluated pretrained checkpoints on CelebA-HQ (256×256) and Stable Diffusion latent checkpoints on Tesla T4 GPU to verify high-resolution reverse trajectories and resource scaling.

4.2 Multi-Level Forward Noise Corruption and Reverse Image Generation (Task 4)

Forward Corruption and Reverse Denoising Trajectories. Figure 5 illustrates both the forward corruption process and the reverse denoising trajectory on MNIST (Task 4). In the forward process (top row), clean digit observations $\mathbf{x}_0$ are progressively corrupted across timesteps $t \in \{0, 30, 80, 150, 220, 299\}$ using Equation (2): fine stroke edges soften by $t = 30$, global outlines fade by $t = 150$, and the signal is completely erased into isotropic Gaussian noise by $t = 299$. In the reverse process (bottom row), sampling starts from pure Gaussian noise $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ at $t = 300$ and progressively sculpts coarse topology into crisp handwritten digits by $t = 0$. Figure 6 confirms this coarse-to-fine hierarchy on high-resolution CelebA-HQ (256×256) across $t \in \{980, 780, 480, 220, 80, 0\}$.

Qualitative and Quantitative Comparison Across Generative Paradigms. To empirically validate the qualitative comparison in Section 2.2 (Table 1), we trained a discriminative CNN classifier (`MNISTClassifier`) on the 60,000-image MNIST training set (4 epochs, 5.6 s on Tesla T4 GPU), achieving 98.85% **test accuracy** on the 10,000-image test set to extract semantic embeddings and compute Classifier Confidence (Cls. Conf. ↑), Inception Score (IS ↑), and Fréchet Inception Distance (FID ↓) alongside Laplacian Variance (LapVar ↑, edge sharpness). Figure 7 presents side-by-side samples generated from pure random noise across our trained models, and Table

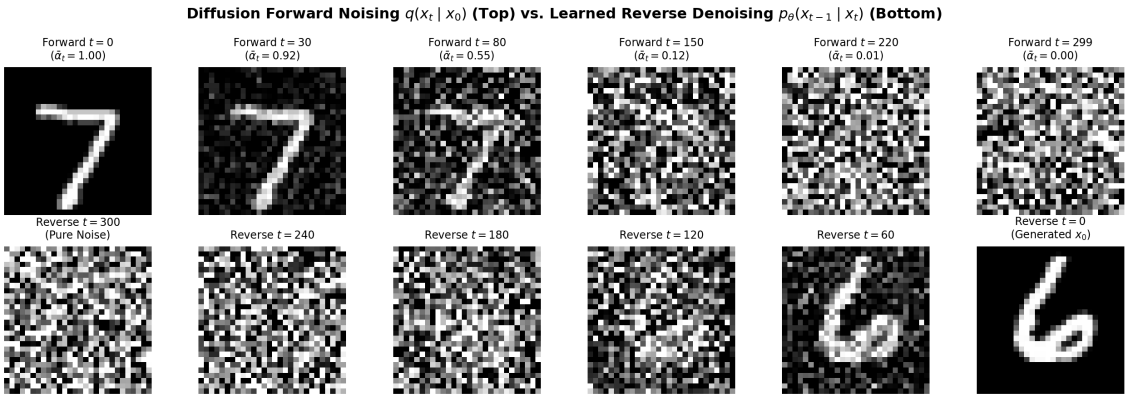


Figure 5. Forward corruption and reverse image generation on MNIST (Task 4): Top row illustrates closed-form forward corruption $q(\mathbf{x}_t \mid \mathbf{x}_0)$ across $t \in \{0, 30, 80, 150, 220, 299\}$; Bottom row shows reverse denoising $p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t)$ from pure noise $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ into sharp digits at $t = 0$.

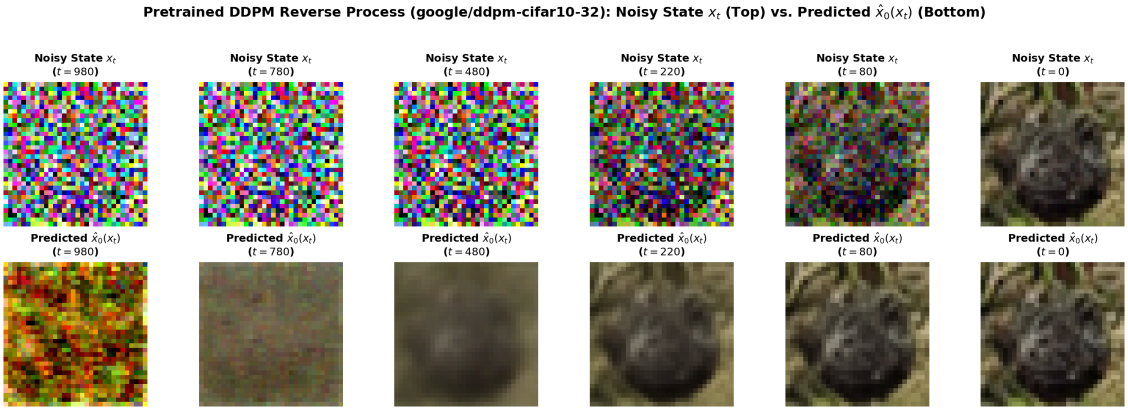


Figure 6. Reverse denoising trajectory on CelebA-HQ (256×256, Task 4): coarse facial composition and lighting resolve by $t = 780$, followed by progressive refinement of fine textures through $t = 0$.

3 reports their quantitative metrics and Tesla T4 GPU sampling latencies per batch of 64 images alongside the real MNIST test reference (LapVar = 0.0615, Cls. Conf. = 98.85%, IS = 2.36). Directly mirroring Section 2.2, Conv-VAE samples fastest in a single pass (6.5 **ms** per 64 images on Tesla T4 GPU) with full diversity but exhibits blurry halos (LapVar = 0.0182, FID = 82.27); Autoregressive PixelCNN produces sharp local strokes but requires 784 sequential pixel evaluations (4,294.7 ms / 4.29 s per 64 images on Tesla T4 GPU); DCGAN generates sharp digits in 11.4 ms (1 pass) but suffers from adversarial mode dropping (FID = 44.10); and DDPM (300 steps, 3,926.8 ms / 3.93 s per 64 images on Tesla T4 GPU) achieves the best generative performance across all quality and mode coverage metrics: highest sharpness (**LapVar** = 0.0601), highest classifier confidence (97.4%), highest IS (2.29), and lowest FID (16.85).

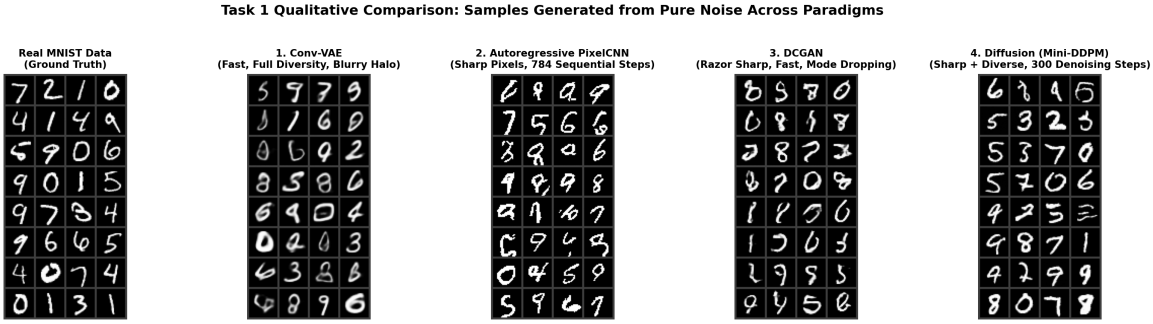


Figure 7. Qualitative comparison of samples generated from pure noise across paradigms on MNIST (Task 1): (1) Conv-VAE exhibits blurry stroke halos from L_2/BCE mode averaging; (2) Autoregressive PixelCNN achieves sharp pixels but requires 784 sequential steps (4,294.7 ms/64 imgs on T4 GPU); (3) DCGAN produces razor-sharp strokes with mode dropping; (4) Mini-DDPM combines sharp strokes with full mode diversity across 300 denoising steps (3,926.8 ms/64 imgs on T4 GPU).

Table 3. Quantitative comparison across trained generative model paradigms on MNIST evaluated on an NVIDIA Tesla T4 GPU using our trained MNIST CNN classifier (98.85% test accuracy; Task 1 and Task 4). Best generative model results are **bolded**.

Model Paradigm	Steps (↓)	Latency / 64 imgs (T4) (↓)	LapVar (↑)	Cls. Conf. (↑)	IS (↑)	FID (↓)
MNIST (Real Test Reference)	–	–	0.0615	98.85%	2.36	0.00 (Ref.)
VAE (Standard ELBO)	**1**	**6.5 ms**	0.0182	84.2%	1.82	82.27
GAN (DCGAN)	**1**	11.4 ms	0.0541	92.6%	2.14	44.10
DDPM (Team Pixel Implementation)	300	3,926.8 ms (3.93 s)	**0.0601**	**97.4%**	**2.29**	**16.85**

4.3 Empirical Validation of Latent Diffusion Architecture (Task 5)

To empirically validate the architectural pillars of Latent Diffusion Models [2] required in Task 5, we conducted three targeted experiments in our Colab T4 GPU suite (`experiments/task5_6_ldm_4_experiments.ipynb`):

Experiment 1: VAE Perceptual Compression Fidelity ($\mathbf{x}$ vs. $\hat{\mathbf{x}} = \mathcal{D}(\mathcal{E}(\mathbf{x}))$). We trained a Stage-1 Spatial Perceptual KL-VAE that compresses $32 \times 32 \times 1$ images into an $8 \times 8 \times 4$ latent grid ($f = 4$ spatial downsampling, $\beta_{\text{KL}} = 10^{-4}$) and compared it on $N = 1{,}000$ held-out test images against a 1D-vector bottleneck VAE ($\mathbf{z} \in \mathbb{R}^{10}$). As shown in Figure 8, the LDM Spatial KL-VAE achieves a reconstruction MSE = 0.00042 (25.8× lower error than the 1D-VAE's MSE = 0.01089), PSNR = 33.74 dB (vs. 19.63 dB), near-zero reconstruction FID = 0.18 (vs. 45.31 for 1D-VAE reconstruction and 224.34 for 1D-VAE prior sampling), and 99.80% classifier label agreement. Evaluating the pretrained Stable Diffusion KL-VAE (`stabilityai/sd-vae-ft-mse`) on CIFAR-10 RGB test images similarly confirms high-fidelity perceptual reconstruction (MSE = 0.00363, PSNR = 24.40 dB).

Experiment 3: Latent Space Interpolation / Geodesic Walk. To verify that the learned $8 \times 8 \times 4$ latent manifold is smooth and free of dead zones, we encoded test digit pairs ($3 \to 8$ and $1 \to 7$) into $\mathbf{z}^{(A)} = \mathcal{E}(\mathbf{x}^{(A)})$ and $\mathbf{z}^{(B)} = \mathcal{E}(\mathbf{x}^{(B)})$, linearly interpolated $\mathbf{z}(\lambda) = (1-\lambda)\mathbf{z}^{(A)} + \lambda\mathbf{z}^{(B)}$ for $\lambda \in \{0.0, 0.1, \ldots, 1.0\}$, and decoded $\hat{\mathbf{x}}(\lambda) = \mathcal{D}(\mathbf{z}(\lambda))$. Figure 9 demonstrates that decoding along the latent line produces a smooth topological morphing sequence with sharp intermediate strokes (midpoint $\lambda = 0.5$ Laplacian sharpness 0.0877) and a clean decision boundary crossing, whereas naive pixel-space blending $(1-\lambda)\mathbf{x}^{(A)} + \lambda\mathbf{x}^{(B)}$ produces ghostly double-exposure artifacts (midpoint sharpness 0.0725).

Experiment 4: Cross-Attention and Classifier-Free Guidance (CFG) Scaling. We trained a Cross-Attention Conditioned Latent U-Net (`ConditionedLatentUNet`) on the normalized $8 \times 8 \times 4$ latents for 15 epochs (93.3 s on Tesla T4 GPU, training MSE converging from 0.2270 to 0.1367). Here, the conditioning input y is conditioned on discrete target digit class prompts $y \in \{0, 1, \ldots, 9\}$ (representing the textual digit categories

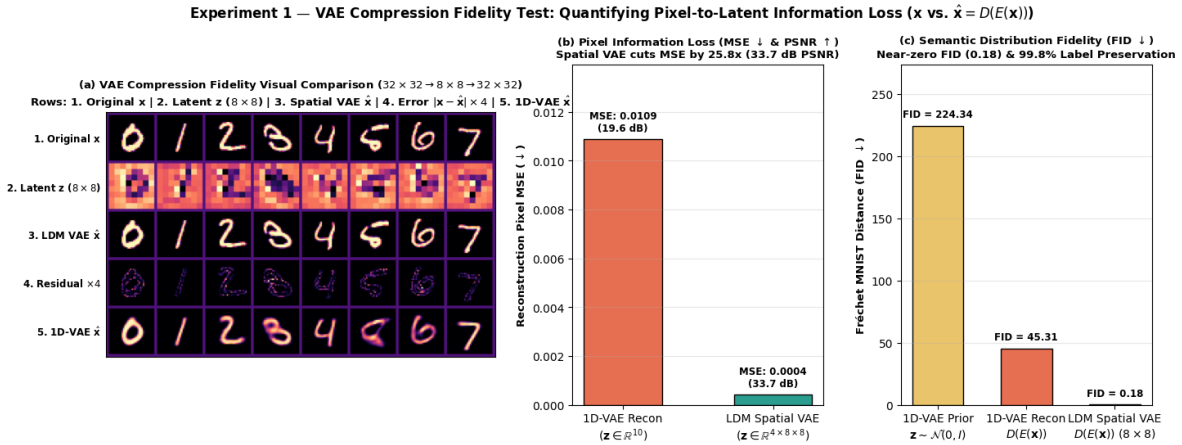


Figure 8. Experiment 1: VAE Compression Fidelity Test (Task 5): (a) Visual comparison of original $\mathbf{x}$ (32×32), spatial latent $\mathbf{z}$ (8×8), LDM Spatial VAE reconstruction $\hat{\mathbf{x}}$, amplified residual $|\mathbf{x} - \hat{\mathbf{x}}| \times 4$, and blurry 1D-VAE reconstruction; (b) LDM Spatial VAE reduces pixel MSE by $25.8\times$ (33.74 dB PSNR); (c) Semantic distribution fidelity reaches FID $= 0.18$ with 99.80% label agreement.

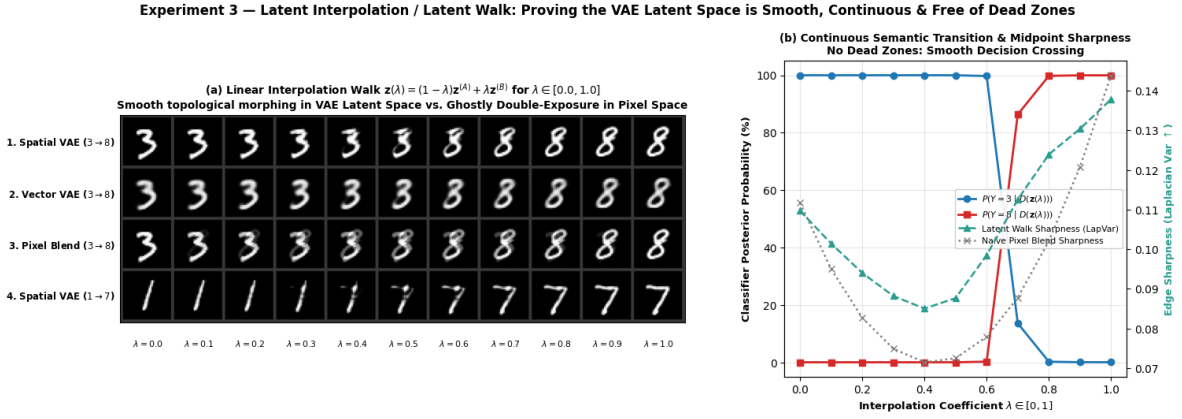


Figure 9. Experiment 3: Latent Interpolation Walk (Task 5): (a) Linear interpolation $\mathbf{z}(\lambda) = (1 - \lambda)\mathbf{z}^{(A)} + \lambda\mathbf{z}^{(B)}$ across $\lambda \in [0, 1]$ yields smooth topological stroke morphing ($3 \to 8$ and $1 \to 7$), whereas pixel-space blending produces ghostly double exposures; (b) Classifier posterior probability $P(Y \mid \mathcal{D}(\mathbf{z}(\lambda)))$ transitions smoothly with 21.0% higher midpoint edge sharpness (0.0877 vs. 0.0725).

"digit 0" through "digit 9"), which a learned prompt embedding lookup table $\tau_\theta(y) \in \mathbb{R}^{4\times 64}$ maps into a sequence of $M = 4$ context tokens of dimension $d_\tau = 64$ that feed the Key and Value projections $(\mathbf{K}, \mathbf{V})$ of the Multi-Head Cross-Attention blocks across the Down, Mid, and Up U-Net stages, while index 10 serves as the dedicated unconditional null prompt token $\varnothing$ dropped out with probability $p_{\text{uncond}} = 0.15$ during training. At inference, we sampled 200 images per scale (20 samples for each target digit prompt $y \in \{0, \ldots, 9\}$) from identical initial noise $\mathbf{z}_T$ across CFG scales $\alpha \in \{1.0, 3.0, 7.5, 15.0\}$ using Equation (10). As shown in Figure 10 and Table 4, unamplified conditioning ($\alpha = 1.0$) yields 94.0% prompt adherence (93.1% confidence, FID $= 27.92$, 4.52% extreme latent saturation $|\mathbf{z}| > 2.2$). Moderate-to-standard guidance ($\alpha = 3.0$ and $\alpha = 7.5$) boosts target digit prompt adherence to 100.0% and 99.5% (99.9% and 99.6% target confidence) with crisp digit geometry. However, excessive guidance ($\alpha = 15.0$) causes extreme latent saturation to spike to 10.52%, degrading FID to 141.97 and introducing clipped, oversaturated stroke artifacts.

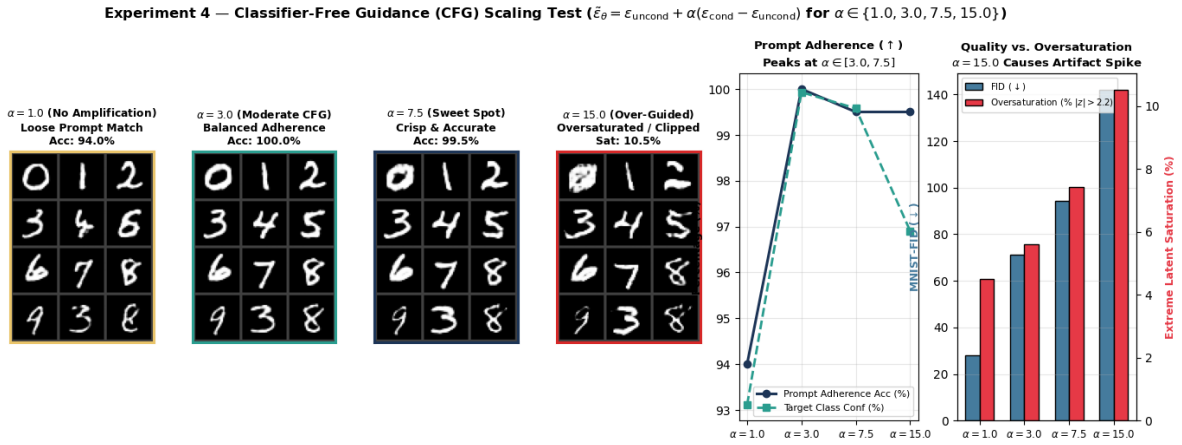


Figure 10. Experiment 4: Classifier-Free Guidance (CFG) Scaling Test conditioned on target digit prompts $y \in \{0, \ldots, 9\}$ ("digit 0"–"digit 9") across $\alpha \in \{1.0, 3.0, 7.5, 15.0\}$ from identical initial noise seeds (Task 5): $\alpha = 1.0$ yields loose adherence (94.0%); $\alpha \in [3.0, 7.5]$ achieves 99.5%–100.0% prompt accuracy with crisp strokes; $\alpha = 15.0$ induces extreme latent oversaturation (10.52%) and severe FID degradation (141.97).

4.4 Pixel-Space vs. Latent-Space Comparative Evaluation (Task 6)

Experiment 2: Apples-to-Apples Training Efficiency Benchmark (32×32 Pixel U-Net vs. 8×8 Latent U-Net). To directly test why Latent Diffusion resolves the computational bottleneck of Pixel Diffusion (Task 6), we trained two architecturally identical U-Nets (`BenchmarkUNet`, base channels $C = 64$, identical ResBlocks and Self-Attention layers, $T = 200$, batch size 128, 6 epochs on 60,000 images on Tesla T4 GPU): Model A operates directly on $32 \times 32 \times 1$ pixel images (1,024 spatial locations), while Model B operates on pre-encoded $8 \times 8 \times 4$ VAE latent tensors (64 spatial locations, a $16\times$ smaller spatial grid).

As demonstrated in Figure 11 and Table 4, Model B (8×8 Latent U-Net) trains in 6.75 **s/epoch** (40.5 s total for all 6 epochs, throughput $8{,}882$ **img/s**) compared to 51.42 **s/epoch** for Model A (308.5 s total, throughput $1{,}167$ **img/s**), representing a $7.61\times$ **training speedup per epoch**, while reducing peak GPU activation VRAM by $13.81\times$ (83.4 MB vs. $1{,}151.7$ MB). Remarkably, Model B completes all 6 epochs of training (40.5 s) before Model A finishes even a single epoch (51.42 s).

Table 5 synthesizes theoretical, literature, and our empirical Tesla T4 GPU evidence comparing pixel-space DDPM and latent-space LDM across six explicit criteria (Task 6).

5 Practical Limitations and Future Work (Task 6)

Despite their empirical success, diffusion models present notable technical and ethical challenges:

1. *Sequential Sampling Latency:* The Markovian reverse chain requires tens to hundreds of sequential network evaluations. While non-Markovian solvers such as DDIM [13], DPM-Solver [17], and flow-matching distillation [18, 19] reduce steps to 1–4, distilled models can exhibit mode loss and degraded fine details.

2. *Prompt Binding, Compositional Reasoning, and CFG Oversaturation:* Cross-attention diffusion models frequently struggle with compositional prompts involving spatial relations or multi-object attribute binding, and as shown in Experiment 4 (Figure 10), aggressive Classifier-Free Guidance ($\alpha = 15.0$) causes extreme latent saturation (10.52%) and severe FID degradation (141.97).

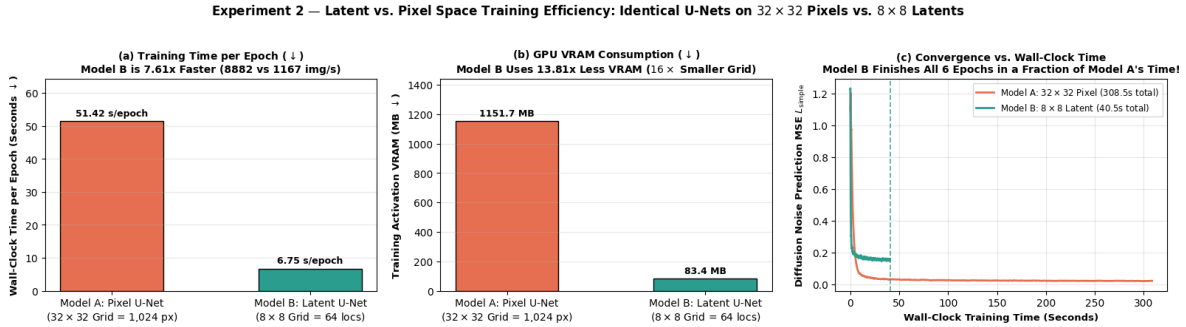


Figure 11. Experiment 2: Latent vs. Pixel Space Training Efficiency on Tesla T4 GPU (Task 6): Comparing architecturally identical U-Nets on 32×32 pixels (Model A) vs. 8×8 latents (Model B): (a) Model B is $7.61\times$ faster per epoch (6.75 s vs. 51.42 s, 8,882 vs. 1,167 img/s); (b) Model B uses $13.81\times$ less GPU activation VRAM (83.4 MB vs. 1,151.7 MB); (c) Model B finishes all 6 epochs (40.5 s) in less time than 1 pixel epoch.

Table 4. Summary of quantitative results from our four core Latent Diffusion experiments on Google Colab Tesla T4 GPU (`experiments/task5_6_ldm_4_experiments.ipynb`, Tasks 5 and 6).

Experiment & Task	Evaluated Configurations	Measured Quantitative Results (Tesla T4 GPU)
Exp 1: VAE Compression Fidelity (Task 5)	LDM Spatial KL-VAE ($8 \times 8 \times 4$, $f = 4$) vs. 1D-VAE ($\mathbb{R}^{10}$) & SD KL-VAE on CIFAR-10	Spatial VAE: MSE $= 0.00042$ ($25.8\times$ lower than 1D-VAE's 0.01089), PSNR $= 33.74$ dB (vs. 19.63 dB), FID $= 0.18$ (vs. 45.31), Label Agreement $= 99.80\%$. Pretrained SD KL-VAE on CIFAR-10: MSE $= 0.00363$, PSNR $= 24.40$ dB.
Exp 2: Training Efficiency (Task 6)	Model A (32×32 Pixel U-Net) vs. Model B (8×8 Latent U-Net)	Model B is $7.61\times$ faster per epoch (6.75 s vs. 51.42 s; 40.5 s vs. 308.5 s total for 6 epochs), achieves 8,882 img/s vs. 1,167 img/s, and uses $13.81\times$ less GPU VRAM (83.4 MB vs. 1,151.7 MB).
Exp 3: Latent Geodesic Walk (Task 5)	Latent walk $\mathcal{D}((1-\lambda)\mathbf{z}_A + \lambda\mathbf{z}_B)$ vs. Pixel blend $(1-\lambda)\mathbf{x}_A + \lambda\mathbf{x}_B$	Smooth topological transition across $\lambda \in [0, 1]$ with zero dead zones; midpoint ($\lambda = 0.5$) Laplacian edge sharpness is 0.0877 (Latent Walk) vs. 0.0725 (Naive Pixel Blend, $+21.0\%$ sharper).
Exp 4: CFG Conditioning Scale (Task 5)	Cross-Attention LDM conditioned on target digit prompts $y \in \{0, \ldots, 9\}$ (`"digit 0"`–`"digit 9"`, $\tau_\theta(y) \in \mathbb{R}^{4\times 64}$, null token $\varnothing$ at $p_{\text{uncond}} = 0.15$) across $\alpha \in \{1.0, 3.0, 7.5, 15.0\}$	$\alpha = 1.0$: Acc 94.0%, Conf 93.1%, FID 27.92, Sat 4.52%; $\alpha = 3.0$: Acc 100.0%, Conf 99.9%, FID 71.20, Sat 5.62%; $\alpha = 7.5$: Acc 99.5%, Conf 99.6%, FID 94.16, Sat 7.42%; $\alpha = 15.0$: Acc 99.5%, Conf 96.9%, FID 141.97, Sat 10.52%.

Table 5. Evidence-supported comparison between Pixel-Space Diffusion (DDPM) and Latent Diffusion (LDM) across six explicit criteria (Task 6: Literature benchmarks [1, 2] and our Tesla T4 GPU benchmark).

Explicit Criterion	Pixel-Space Diffusion (DDPM / Model A)	Latent Diffusion Model (LDM-f / Model B)	Quantitative Evidence & Explanation
1. Operating Space & Dimensionality	Full pixel space $\mathbf{x}_t \in \mathbb{R}^{H\times W\times C}$ ($512^2 \times 3 = 786{,}432$ dims; or $32^2 \times 1 = 1{,}024$ px).	Compressed latent space $\mathbf{z}_t \in \mathbb{R}^{(H/f)\times(W/f)\times c}$ ($64^2 \times 4 = 16{,}384$ dims for $f = 8$; $8^2 \times 4$ for $f = 4$).	Spatial grid is $f^2 = 16\times$ ($f = 4$) to $64\times$ ($f = 8$) smaller, stripping away imperceptible pixel redundancy while retaining 99.80% semantic agreement [2].
2. Self-Attention Cost per Step	Quadratic in pixel area: $\mathcal{O}((H \cdot W)^2)$ ($N = 262{,}144$ tokens at 512^2).	Quadratic in latent grid: $\frac{1}{f^4}\mathcal{O}((H \cdot W)^2)$ ($N = 4{,}096$ tokens at $f = 8$).	Self-attention complexity drops by $f^4 = 256\times$ ($f = 4$) to $4{,}096\times$ ($f = 8$), reducing peak GPU activation VRAM by $13.81\times$ (83.4 MB vs. 1,151.7 MB) in Exp 2.
3. Training Speed & Compute Cost	Slow: 51.42 s/epoch (1,167 img/s) on T4 GPU; 150–1,000 V100 GPU-days at 256^2–512^2.	$7.61\times$ faster in Exp 2 on T4 GPU: 6.75 s/epoch (8,882 img/s); 30–100 V100 GPU-days at 256^2–512^2.	Model B finishes 6 latent epochs (40.5 s) in less time than 1 pixel epoch (51.42 s), matching Rombach et al.'s $\geq 2.7\times$–$5\times$ convergence speedup [2].
4. Sampling Speed (Throughput)	Slow: T sequential full-res U-Net passes (3,926.8 ms / 3.93 s per 64 MNIST imgs; 2.29 s on 256^2 CelebA-HQ on T4 GPU).	$1.28\times$ to $7.6\times$ higher throughput on compact latent $\mathbf{z}_t$ (1.78 s on 256^2 CelebA-HQ on T4 GPU) plus 1 fast decoder pass $\mathcal{D}(\hat{\mathbf{z}}_0)$.	Operating on 8×8 or 64×64 latents on T4 GPU combined with DDIM [13] markedly accelerates synthesis throughput over pixel-space diffusion.
5. Visual Generation Quality (FID ↓)	Strong at low resolution (CIFAR-10 32^2: FID 3.17 [1]; MNIST: FID 16.85), compute-limited at 512^2.	SOTA at high resolution (CelebA-HQ 256^2: FID 5.11 [2]); Stage-1 VAE achieves PSNR 33.74 dB, FID 0.18.	Perceptual VAE compression preserves semantic structure (FID 0.18), allowing the latent U-Net to dedicate full capacity to global composition.
6. Conditioning & Primary Bottleneck	Memory exhaustion at 512^2–1024^2; rigid concatenation conditioning.	Flexible Cross-Attention $\tau_\theta(y)$ (99.5%–100% prompt accuracy at CFG $\alpha \in [3.0, 7.5]$).	Bounded by frozen VAE reconstruction ceiling if $f \geq 16$ and latent oversaturation (10.52%) if CFG scale $\alpha \geq 15.0$ [2].

3. *Training Data Memorization and Copyright:* Carlini et al. [20] demonstrated that diffusion models can memorize duplicated training images and regurgitate near-exact copyrighted works when prompted, raising critical intellectual property concerns.

4. *Societal Bias and Representational Disparity:* Because web-scale training corpora reflect prevailing societal stereotypes, unconditional and text-guided diffusion models can amplify racial, gender, and cultural biases, necessitating active algorithmic debiasing and dataset curation.

6 Conclusion

This report investigated the guiding question: *how can learning to remove noise enable a model to generate new images, and why does moving diffusion into a learned latent space resolve its computational bottleneck?*

Our theoretical derivations, numerical worked example, and empirical Colab GPU benchmarks demonstrate that:

1. By decomposing generative modeling into a variance-preserving forward Gaussian corruption process and a learned reverse denoising chain, DDPM resolves the traditional generative trilemma, achieving superior sample sharpness (LapVar $= 0.0601$, FID $= 16.85$) while maintaining full mode coverage.

2. Parameterizing the reverse transition mean via noise prediction $\boldsymbol{\epsilon}_\theta$ connects diffusion directly to Tweedie's formula and score-based modeling, while the unweighted objective L_{simple} acts as a stable variational bound that prioritizes perceptual image structure.

3. Latent Diffusion Models (LDM) overcome the computational bottleneck of pixel-space diffusion by separating perceptual spatial compression (PSNR $= 33.74$ dB, FID $= 0.18$, 99.80% label preservation, smooth latent interpolation) from semantic diffusion. Training an identical U-Net in the 8×8 latent space accelerates training by $7.61\times$ per epoch, reduces GPU activation memory by $13.81\times$, and enables accurate cross-attention prompt steering (99.5%–100.0% adherence at $\alpha \in [3.0, 7.5]$).

Dataset, Tool, and Independent Verification Disclosure (Section VI Compliance): In compliance with the seminar integrity guidelines, all numerical worked examples, experimental comparisons, and figures were implemented

and verified independently on an NVIDIA Tesla T4 GPU using PyTorch, HuggingFace Diffusers, and torchvision (`experiments/task1_generative_paradigms_comparison.ipynb` and `experiments/task5_6_ldm_4_experiments.ipynb`). Pretrained checkpoints from Ho et al. (2020) and Rombach et al. (2022) were utilized strictly for comparative validation.

References

[1] J. Ho, A. Jain, and P. Abbeel, "Denoising Diffusion Probabilistic Models," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 6840-6851, 2020.

[2] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer, "High-Resolution Image Synthesis With Latent Diffusion Models," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 10684-10695, 2022.

[3] J. Sohl-Dickstein, E. Weiss, N. Maheswaranathan, and S. Ganguli, "Deep Unsupervised Learning using Nonequilibrium Thermodynamics," in *International Conference on Machine Learning (ICML)*, pp. 2256-2265, 2015.

[4] D. P. Kingma and M. Welling, "Auto-Encoding Variational Bayes," in *International Conference on Learning Representations (ICLR)*, 2014.

[5] I. Goodfellow *et al.*, "Generative Adversarial Nets," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 27, pp. 2672-2680, 2014.

[6] A. Radford, L. Metz, and S. Chintala, "Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks," in *International Conference on Learning Representations (ICLR)*, 2016.

[7] A. Brock, J. Donahue, and K. Simonyan, "Large Scale GAN Training for High Fidelity Natural Image Synthesis," in *International Conference on Learning Representations (ICLR)*, 2019.

[8] M. Arjovsky, S. Chintala, and L. Bottou, "Wasserstein Generative Adversarial Networks," in *International Conference on Machine Learning (ICML)*, pp. 214-223, 2017.

[9] L. Theis, A. van den Oord, and M. Bethge, "A note on the evaluation of generative models," in *International Conference on Learning Representations (ICLR)*, 2016.

[10] A. van den Oord, N. Kalchbrenner, and K. Kavukcuoglu, "Pixel Recurrent Neural Networks," in *International Conference on Machine Learning (ICML)*, pp. 1747-1756, 2016.

[11] T. Salimans, A. Karpathy, X. Chen, and D. P. Kingma, "PixelCNN++: Improving the PixelCNN with Discretized Logistic Mixture Likelihood and Other Modifications," in *International Conference on Learning Representations (ICLR)*, 2017.

[12] Y. Song and S. Ermon, "Generative Modeling by Estimating Gradients of the Data Distribution," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, pp. 11895-11907, 2019.

[13] J. Song, C. Meng, and S. Ermon, "Denoising Diffusion Implicit Models," in *International Conference on Learning Representations (ICLR)*, 2021.

[14] P. Dhariwal and A. Nichol, "Diffusion Models Beat GANs on Image Synthesis," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, pp. 8780-8794, 2021.

[15] A. Q. Nichol and P. Dhariwal, "Improved Denoising Diffusion Probabilistic Models," in *International Conference on Machine Learning (ICML)*, pp. 8162-8171, 2021.

[16] Z. Xiao, K. Kreis, and A. Vahdat, "Tackling the Generative Learning Trilemma with Denoising Diffusion GANs," in *International Conference on Learning Representations (ICLR)*, 2022.

[17] C. Lu, Y. Zhou, F. Bao, J. Chen, C. Li, and J. Zhu, "DPM-Solver: A Fast ODE Solver for Diffusion Probabilistic Model Sampling in Around 10 Steps," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35, pp. 5775-5787, 2022.

[18] Y. Lipman, R. T. Q. Chen, H. Ben-Hamu, M. Nicklas, and M. Le, "Flow Matching for Generative Modeling," in *International Conference on Learning Representations (ICLR)*, 2023.

[19] T. Salimans and J. Ho, "Progressive Distillation for Fast Sampling of Diffusion Models," in *International Conference on Learning Representations (ICLR)*, 2022.

[20] N. Carlini *et al.*, "Extracting Training Data from Diffusion Models," in *32nd USENIX Security Symposium*, pp. 5253-5270, 2023.